Chapter 3

Learning and Evaluation

Xu Kuang, Stanford GSB

We have seen in the earlier chapter how to use stochastic modeling (the Erlang M/M/c/k systems) and the resulting Erlang formulas to guide decisions around how many processing resources to allocate for a given level of throughput and probability of delay. The discussion there, however, was largely predicated on us accepting the model as already determined and readily available, and also that it correctly captures reality.

Recall from Chapter 1 the Frame-Learn-Use-Evaluate modeling workflow. In this chapter, we will discuss the procedure for learning a specific congestion model within the broader family and the questions that arise concerning the eventual effectiveness of this learned model in guiding actions. The procedure for learning such a model is typically neither very difficult nor complex, but evaluating whether we have ended up with a good enough model can be more subtle and delicate.

3.1 The Basics of Learning the Congestion Model

In the case of congestion modeling in the Erlang style, we implicitly, via human expertise, frame the problem by adopting a mechanistic stochastic model that depicts the congestion dynamics arising from serving randomly arriving requests using the family of M/M/c/k queueing models. Each individual model within this family is further specified by several parameters. By learning a congestion model, we mean identifying the most appropriate combination of these parameters. They include:

1. The arrival rate λ: how many jobs arrive in the system per unit of time?
2. The service rate μ: how many jobs can each server or worker process per unit of time?
3. The number of servers or workers c.
4. The total number of jobs the system can contain, k. Alternatively, the size of the waiting buffer, $k - c$.

Out of these parameters, c and k are usually not considered something to be learned because they can be chosen by the designer themselves. This leaves determining the values of λ and μ to be learned from data.

Cold- vs. Warm-starts

In the case of a call center or data center, λ and μ are typically estimated using operational data from existing operations. The emphasis here is that we are assuming that *some* form of the system the model aims to capture has already been deployed in practice for long enough to collect such data. We call this the *warm-start* modeling problem, where the system is already in operation, affording modelers like Erlang some readily available data to get started. In contrast, there are equally many applications, if not more prevalent, in which the decision-maker faces a *cold-start* problem, where we hope to use modeling

to guide the design of a system that has yet to be launched or has only recently been launched.

The detailed statistical procedures employed to estimate and learn these parameters from data can vary greatly depending on the preference of the designer (Bayesian, Frequentist, etc.). As a baseline for our example, things are really quite simple. By choosing such a relatively simple queueing model, we are effectively baking in the assumption that the arrival and service rates are time-invariant, and in particular, that they do not depend on the day of week or time of day. With this restriction in place, and suppose that we are learning the parameter for a model of a call center, learning λ simply amounts to:

1. Define a sufficiently long data collection horizon, T. Say we let this be one year, so $T = 525,600$ minutes.

2. Measure the number of incoming calls during this year, N. Suppose this number came to be $N = 28,598,331$ calls.

3. Then λ is estimated to be $N/T = 54.41$ calls per minute.

The estimation of μ can be accomplished in a similar manner, by taking the average of processing times of jobs over a sufficiently long time horizon.

To be more precise, the estimation procedure above corresponds to the frequentist maximum-likelihood estimator for λ underz the time-invariance assumption. Some practitioners may also prefer a Bayesian procedure, where they can incorporate stronger prior beliefs about what λ should be. In this example, a large number of observations (across a whole year) is collected to estimate a relatively large arrival rate, and as such these two approaches are unlikely to yield different conclusions. But in cases where arrivals occur sparingly, such as when estimating dynamics related to highly specialized equipment in the METRIC model discussed in Chapter 2, the practitioner's prior could play a much more significant role.

3.2 Wrong Input, Right-ish Output: Do Erlang Models Capture Reality?

Being able to estimate the parameters in a congestion model is, however, only half the battle. A common mode of failure of stochastic modeling in the real world stems from choosing the wrong model family to begin with. That is, implicit assumptions deep within the constitution of the model can make it so fundamentally deviant from reality from the outset that even sound statistical procedures for estimating the various parameters cannot make up for these more fundamental departures.

In the Erlang examples, there are many such possible assumptions that could, in theory, prevent the model from reliably portraying reality. To name just a few, inter-arrival times may not be exponentially distributed (a fact called out by Erlang himself in the original 1909 paper). The arrival rate is very likely not stationary over time in almost any real-world application. Servers may not work at the same speed. And finally, customers may not patiently wait in line to receive service indefinitely. Some of these assumptions may be necessary evils to cope with in exchange for a relatively elegant and robust mathematical model that is still useful enough, while others may be so egregious as to become deal-breakers. Without further evaluation, a practitioner simply cannot tell one from another.

Brown et al. (2005) offer an interesting case study for evaluating Erlang-style models against real-world operating conditions. The study involves call-by-call operational data from a small bank call center in Israel during the full year of 1999. Within the 1.2 million incoming calls in the dataset, about 750,000 terminate in an automated Voice Responder Unit (VRU) and do not reach any human agents, while the remaining 450,000 requests go past the VRU before queueing for a human agent. Even here, before any analysis begins, we see that the Erlang model may capture only part of a larger real-world system, given that more than 60% of the calls did not even enter this queueing system to begin with. The paper then examines a number of common modeling assumptions against the empirical data. Some of the findings are as one would expect, while others are quite surprising.

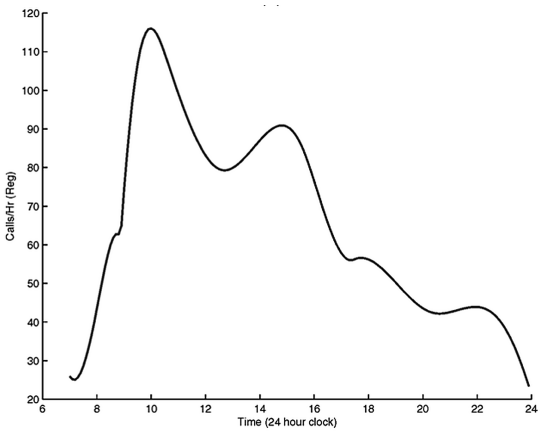

(a) Arrival-rate pattern (Figure 1).

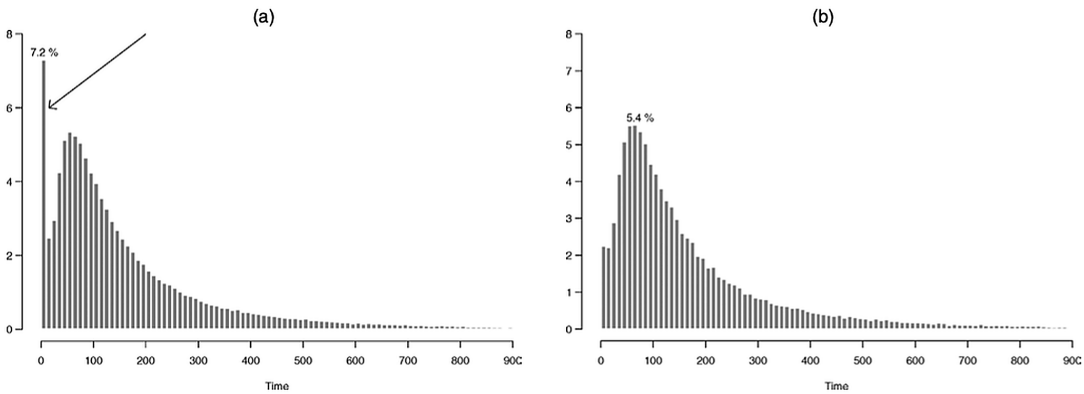

(b) Service-time distribution (Figure 2).

Figure 3.1: Empirical diagnostics from (Brown et al., 2005): (a) arrival-rate pattern and (b) service-time distribution.

For starters, two core assumptions in Erlang models are that (a) calls arrive at a constant rate according to a Poisson process, and (b) the time it takes to process each call is exponentially distributed. The data shows that neither assumption matches reality. Figure 3.1a shows the arrival rate over time, and as one would expect, the rate varies greatly over the course of a day. Figure 3.1b shows the service-time distributions over two different months. In both cases, the distribution is not Exponential. For instance, the

distribution peaks around 80, whereas an Exponential distribution would have the highest probability density around 0. Upon reflection, this is not too surprising, as one might expect that a phone call to a bank would have some "typical" length that is long enough to resolve an issue. In fact, the paper shows that the empirical service-time distribution is much better approximated by a Lognormal distribution, which tends to have a much heavier tail than an Exponential distribution (i.e., there is a greater chance of very long calls)[1].

Having such a mismatch between the core modeling assumptions and empirical data does not bode well for the validity of Erlang models. For non-stationary arrival rates, the issue is perhaps less severe. Practitioners have long circumvented this problem by simply using the Erlang formula in smaller time windows, typically a 30-minute to 1-hour window. It is believed that during these smaller windows, arrival rates are stable enough to be treated as constant. In other words, to model the congestion level over, say, 12 hours, we would run $12 \times 2 = 24$ separate Erlang formulas. Using the Erlang formula with a non-Exponential service-time distribution is, however, a further stretch, because we lose all the convenience that comes with working with a Markov process, and we now know that congestion dynamics depend on the coefficient of variation of the service-time distribution (the ratio between variance and mean squared), which is very different under the assumed Exponential versus the more realistic Lognormal distribution.[2]

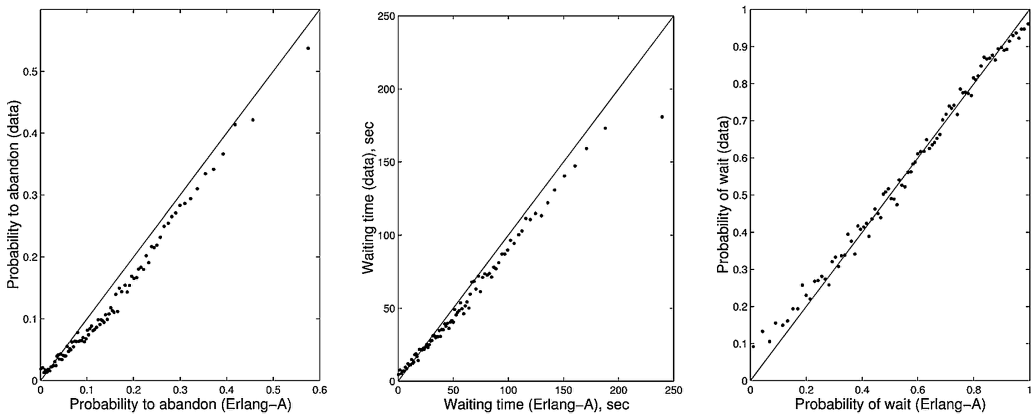

Figure 3.2: Figure 8 in (Brown et al., 2005), comparing modified Erlang-model predictions against observed congestion outcomes.

Given all this mismatch between assumptions and data, one might expect Erlang models to perform poorly at predicting observed congestion metrics. Well... they do not! The researchers then evaluated a modified Erlang model[3] against actual congestion measurements, such as probability of waiting and average waiting times. In Figure 3.2, a total of 3,867 hour-long intervals are partitioned into 45 groups, where each group contains hours with similar congestion loads (arrival rate divided by total service rate), and each dot represents the average measurement from that group. We see that, barring

[1]For a plausible explanation of lognormality via theories of mathematical psychology, see (Ulrich and Miller, 1993).

[2]See (Whitt, 1993).

[3]They added an additional feature to model customer abandonments while waiting in queues, but other than that the model still closely tracks the original Erlang version.

some underestimation of the abandonment probabilities, the Erlang model is generally quite remarkable at producing congestion estimates that closely track empirically observed outcomes. Granted, there is the caveat that each dot is still only an average over about 80 hours, and as such the predictive error per hour may be worse. Even with that caveat, the contrast is striking: despite being built on evidently idealized and "unreal" assumptions, Erlang models still seem to have captured a certain "core essence" of the system dynamics and are able to deliver accurate predictions.

3.3 Recommended Reading

For primary empirical evidence on whether Erlang-style assumptions hold in a modern service system, see (Brown et al., 2005). For original historical context behind the Erlang family itself, see Chapter 2 and (Erlang, 1917).

Bibliography

Amazon Web Services (2024). Amazon Bedrock Service Level Agreement. AWS SLA page. Accessed 2024.

Baseten (2026). How Hebbia Uses Baseten to Power AI Workflows. Baseten customer story. Discusses TTFT-sensitive inference workloads; accessed April 10, 2026.

Bhat, U. N. (1969). Sixty years of queueing theory. *Management Science*, 15(6):B280–B294. Application Series.

Brown, L., Gans, N., Mandelbaum, A., Sakov, A., Shen, H., Zeltyn, S., and Zhao, L. (2005). Statistical analysis of a telephone call center: A queueing-science perspective. *Journal of the American Statistical Association*, 100(469):36–50.

Danmarks Statistik (1901). Folkemængden 1. februar 1901 i Kongeriget Danmark efter de vigtigste administrative inddelinger. Danmarks Statistik publication page. Official census publication; accessed April 2, 2026.

Encyclopaedia Britannica (2026). Loading coil. Britannica article page. Overview of loading-coil purpose in long-distance telephony; accessed April 2, 2026.

Erlang, A. K. (1909). Sandsynlighedsregning og telefonsamtaler. Digital scan (Royal Danish Library). [Digital scan; accessed 2026-03-29].

Erlang, A. K. (1917). Solution of some problems in the theory of probabilities of significance in automatic telephone exchanges. *Elektroteknikeren*, 13:5–13. Original Danish title: Łosning af nogle Problemer fra Sandsynlighedsregningen af Betydning for de automatiske Telefoncentraler. Archived scan: Wayback copy.

Fischer, C. S. (1978). The spatial diffusion of information and the long-distance telephone, 1875–1900. *Social Science History*, 2(2):173–196.

Gingerich, O. (2007). The copernican myths. https://physicstoday.aip.org/features/the-copernican-myths. [Physics Today; accessed 2026-03-22].

Gross, D., Shortle, J. F., Thompson, J. M., and Harris, C. M. (2008). *Fundamentals of Queueing Theory*. Wiley, Hoboken, NJ, 4 edition.

Hajek, B. (2015). *Random processes for engineers*. Cambridge university press. https://hajek.ece.illinois.edu/ECE534Notes.html.

Hebbia (2025). The Distributed System Behind Hebbia's High-Scale AI. Hebbia engineering blog. Published March 17, 2025; accessed April 10, 2026.

INFORMS (2026). Who Was Agner Krarup Erlang? INFORMS Erlang page. INFORMS Erlang Prize background page; accessed April 2, 2026.

Jewett, F. B. (1951). The evolution of inductive loading. *Bell System Technical Journal*, 30(2):123–140.

Kendall, D. G. (1953). Stochastic processes occurring in the theory of queues and their analysis by the method of the imbedded markov chain. *The Annals of Mathematical Statistics*, 24(3):338–354.

Kepler, J. (1609). *Astronomia Nova*. Heidelberg. Digitized edition metadata via e-rara.

Kuhn, T. S. (1957). *The Copernican Revolution: Planetary Astronomy in the Development of Western Thought*. Harvard University Press, Cambridge, MA.

National Park Service (2005). New Jersey Bell Headquarters Building, National Register of Historic Places Registration Form. NPS/NJ historic register PDF. Includes historical summary of the Independent Telephone Movement and entry scale; accessed April 2, 2026.

New York Public Library (2026). 1964–1965 new york world's fair corporation records, 1939–1966. NYPL archives entry. Photo archive; accessed 2026-03-30.

Newton, I. (1687). *Philosophiae Naturalis Principia Mathematica*. Jussu Societatis Regiae ac Typis Josephi Streater, London.

Newton, I. (1999). *The Principia: Mathematical Principles of Natural Philosophy*. University of California Press, Berkeley.

Our World in Data (2026a). Mean years of schooling (long-run). OWID dataset link. Long-run education dataset; accessed April 2, 2026.

Our World in Data (2026b). Share of the population with tertiary education. OWID dataset link. Dataset adapted from Barro and Lee (2015) and Lee and Lee (2016); accessed April 2, 2026.

Ross, S. M. (2014). *Introduction to Probability Models*. Academic Press, Amsterdam, 11 edition.

Shalev-Shwartz, S. and Ben-David, S. (2014). *Understanding machine learning: From theory to algorithms*. Cambridge university press.

Sherbrooke, C. C. (1968). Metric: A multi-echelon technique for recoverable item control. *Operations Research*, 16(1):122–141.

Sherbrooke, C. C. (2004). *Optimal Inventory Modeling of Systems: Multi-Echelon Techniques*. Kluwer Academic Publishers, Boston, 2 edition.

Skee, J. (2019). The history of the analysis of queues and waiting lines. INFORMS history page. Edited by Linus Schrage; INFORMS History of O.R. Excellence, Queueing Models; accessed 2026-03-30.

Skee, J. D. (2016). *By the Numbers: Confidence, Consultants, and the Construction of Mass Leisure, 1953–1975*. Phd dissertation, University of California, Berkeley.

Strevens, M. (2020). *The Knowledge Machine: How Irrationality Created Modern Science*. Liveright Publishing Corporation, New York.

Sutton, R. (2019). The bitter lesson. http://www.incompleteideas.net/IncIdeas/BitterLesson.html. [Online essay; accessed 2026-03-26].

Ulrich, R. and Miller, J. (1993). Information processing models generating lognormally distributed reaction times. *Journal of Mathematical Psychology*, 37:513–525.

U.S. Bureau of Labor Statistics (2026). Cpi inflation calculator. https://www.bls.gov/data/inflation_calculator.htm. [CPI-U based; accessed 2026-04-02].

U.S. Census Bureau (2023). Facts by Decade: 1900. U.S. Census page. Includes 1900 urban population table; accessed April 2, 2026.

Whitt, W. (1993). Approximations for the gi/g/m queue. *Production and Operations Management*, 2:114–161.